

DB101 – Logging and recovery

Goetz Graefe – Madison, Wis.

“ACID” transactions


- ⇒ **A**tomicity: “all or nothing” success or failure
- ⇒ **C**onsistency: logical & physical, e.g., integrity constraints & indexes
- ⇒ **I**solation: concurrency control “equivalent to serial execution”
- ⇒ **D**urability: persistence “come fire, flood, or insurrection”

Many forms of “cheating”:

- reads and writes at different times
- commit sequence \neq commit log record sequence
 - ⇒ no snapshot isolation transactions, no point-in-time recovery
- ...

Update propagation

...must be **reliable**, **simple**, **efficient**, **scalable**, and **robust**!



Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Omitted topics

- Comprehensive consistency checks of b-trees, tables, databases, backups, recovery logs, log archives
- Recovery from loss of recovery log or log archive
- Allocation-only logging
- Logging and recovery of offline and online index operations (creation)
- Logging in two-phase commit, 3pc, Paxos, Raft, etc.
- Compression of recovery log and log archive
- Transaction rollback and 2pc in multi-version indexes
- Microsoft “accelerated database recovery”
- Instant recovery for file systems and database blobs

Agenda

- **Write-ahead logging**
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Write-ahead logging

Recovery log = recorded history of database updates on “stable storage”

No revising history \Rightarrow append-only log storage

- Ever-increasing LSNs (log sequence number = address of log record)
- Sequential writing – no gaps in history on stable storage

Database history: create, alter, drop, insert, update, delete...

Server history: begin transaction, commit transaction, checkpoint...

Transaction commit: server history \Rightarrow database history

Example recovery logs

Transactional database changes
+ server events

Begin transaction
Update page 7, slot 9: ...
Update page 4, slot 6: ...
Rollback
Rollback
Commit (nothing)

T₁: Begin transaction
T₁: Update page 7, slot 9: ...
Database checkpoint
T₁: Update page 4, slot 6: ...
T₂: Begin transaction
T₂: Update page 3, slot 1: ...
Written page 7
T₁: Rollback
T₂: Update page 4, slot 6: ...
T₁: Rollback
T₁: Commit (nothing)

Physical & logical logging

1970s: physical logging = before & after page images

1980s: logical logging: SQL commands & commit sequence

1990s: “physiological” logging: physical “redo”, logical “undo”

- required for row-level locking, e.g., key-value locking in b-trees
- page id + slot #, not byte offset within pages

2000s: logical logging + serial execution in VoltDB

2010s: mirror-don't-log, append-only data stores, log-only storage...
(is a mirror a page log?)


Write-ahead logging

No update-in-place until the “undo” log record is on stable storage

“Redo-undo” log record: new & old values

Allocation-only logging

“Redo-only” log records: rollback actions, system transactions



User transactions & system transactions

User transactions: retrieve & update logical database contents

- Updates observable with “select count(*) from... where...”

System transactions: modify physical database representation

- E.g., b-tree node split, ghost removal
- No locks on database contents, only latches on data structures
- No log flush on commit, single log record \Rightarrow no “undo”
- Invoked within thread, no transaction descriptor or identifier
- *All space management actions!*

User transactions & system transactions

	User transactions	System transactions
Invocation source	User requests	System-internal logic
Database effects	Logical database contents	Physical storage structure
Data location	Database or buffer pool	In-memory page images
Parallelism	Multiple threads possible	Single thread
Invocation overhead	New thread	Same thread
Locks	Acquire and retain	Test for conflicts
Commit overhead	Force log to stable storage	No forcing
Logging	Full “redo” and “undo”	Omit “undo” in most cases
Recovery	Backward	Forward or backward

Ghost records, ghost space, ghost keys...

Transaction 1:

1. Begin transaction
2. Delete from Table1 where ...
3. ...
4. **Abort transaction**

Transaction 2:

1. Begin transaction
2. ...
3. Insert into Table1 values (...)
4. ...

- Space management (free bytes)
- Integrity constraints (uniqueness)

Ghost indexes...

Transaction 1:

1. Begin transaction
2. Drop index Orders.CustIdx
3. ...
4. Abort transaction

Transaction 2:

1. Begin transaction
2. ...
3. Create index PartIdx on LineItem (PartNo...)
4. ...

- Space management (free storage)
- Logging volume (rollback “drop”)

Agenda

- Write-ahead logging
 - **Commit**, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Transaction commit

Single time & place for commit decision

- e.g., user commits herself, then system attempts to commit
- updates become durable and then visible
- force commit log record to recovery log on “stable storage”
transaction commit: server history \Rightarrow database history
- “Force” policy (i.e., flush dirty pages to database) not required!

Latency vs bandwidth – “group commit”

Small transaction $\Rightarrow \leq 1$ KB of log records

Fast device $\Rightarrow 0.1 \text{ ms} \times 200 \text{ MB/s} = 20 \text{ KB}$ page size
(latency \times bandwidth)

Log flush per commit $\Rightarrow >90\%$ of log space & bandwidth wasted!

\Rightarrow group commit = short timeout before log flush, e.g., 1 ms

Two-phase commit

Single time & place for commit decision

- Multiple logs, single commit decision
- Participants “vote”: pledge to realize coordinator’s commit decision
- Pledge must survive lock contention & deadlock, server restart, etc.
- Coordinator decides & logs its decision
- Participants receive & log the decision

Paxos vs two-phase commit – different problems!

Paxos

- Majority agreements on coordinator + outcome
- Orthogonal to concurrency control
- Maintenance of replicas & mirrors

Two-phase commit

- Unanimous pledge, single decision
- Implications for concurrency control (locking)
- Maintenance of shards & partitions

Agenda

- Write-ahead logging
 - Commit, **backups**, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Classification of database backups ~~server backups~~

Contents & frequencies

- Full: all allocated database pages, e.g., every Sunday 2am
- Differential: pages changed since full backup, e.g., since Sunday
- Incremental: pages changed since last backup, e.g., since yesterday

Consistency

- Transaction-consistent: quiescent database server
- Snapshot (point-in-time): transaction-consistent or not
- Fuzzy: database pages as found in buffer & disk-order scan...

Compressed, encrypted, sorted on database page identifier, indexed...

Agenda

- Write-ahead logging
 - Commit, backups, **archives**, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions


Recovery log vs log archive

Recovery log – written during transaction processing

- Low-latency reliable storage, e.g., non-volatile memory or SSDs

Log archive – written asynchronously

- Cheap but very reliable storage, e.g., RAID-5 or RAID-6



Recovery log vs log archive

Recovery log – written during transaction processing

- Low-latency reliable storage, e.g., SSDs or non-volatile memory
 - Chosen for latency, not for capacity or price/capacity
 - No compression, encryption, sorting, indexing...
- Sequence of log records strictly by time

Log archive – written asynchronously

- Chosen for cheap reliable storage, e.g., RAID-5 or RAID-6
 - Usually compressed & encrypted
- Could be sorted & indexed, e.g., on database page identifier

Agenda

- Write-ahead logging
 - Commit, backups, archives, **checkpoints**
- Classic failures & recovery techniques
 - Transaction, system, media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

~~Database checkpoints~~ Server checkpoints

Checkpoint = log record capturing current server state

- Active transactions & their locks
- Dirty pages & their oldest change

1970s checkpoint: quiescent transactions, flushing all dirty pages

1990s checkpoint: flush dirty pages, “second chance”

2010s checkpoint: multiple log records, partitioned checkpoints

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- **Classic failures & recovery techniques**
 - Transaction, system, media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Classic failures & recovery techniques

Classification of all failures into transaction, system, and media failures

- Missing: double failures, node failures, log failures...

Precisely three recovery procedures

- Transaction rollback
- System restart
- Media restore

Scale → ↓ History	Focused	Database-wide
Repeat		“Redo” log scan
Rewind	Transaction rollback	“Undo” log scan

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - **Transaction**, system, media **failures**; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Transaction failure \Rightarrow transaction rollback

Any cause: deadlock, timeout, user cancel, global abort decision...

Linked list of log records per transaction:

- LastLSN in transaction descriptor (in-memory transaction manager)
- PriorLSNsameTx in each log record

1970s: idempotent “undo” – no log of rollback

1990s: “compensation log records” in ARIES

2010s: “rollback log records”

Transaction failure \Rightarrow transaction rollback

1970s: idempotent “undo” – no log of rollback \Leftarrow restart with full log

- “Abort” log record, no end-of-transaction log flush

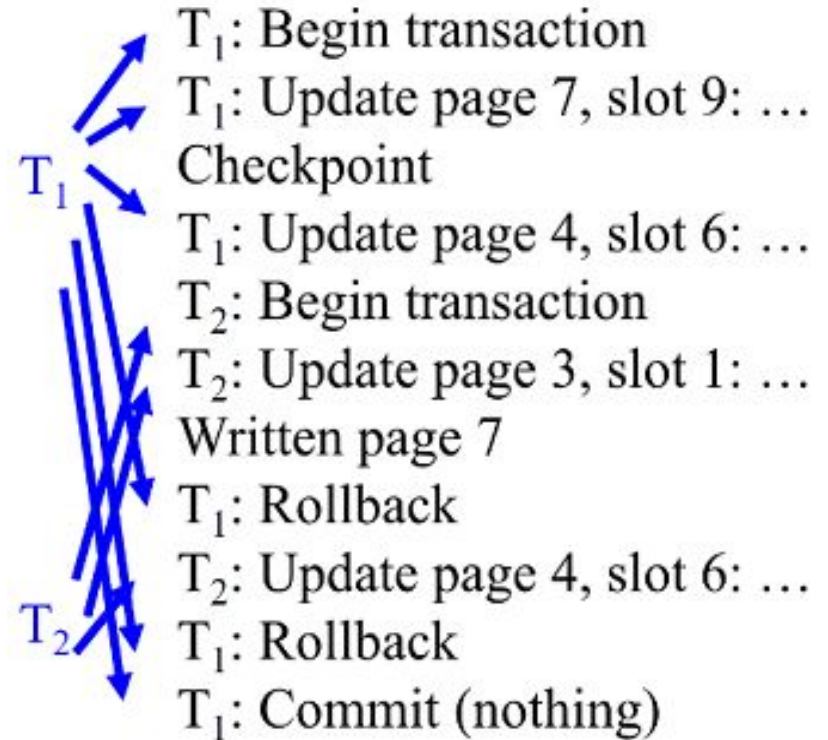
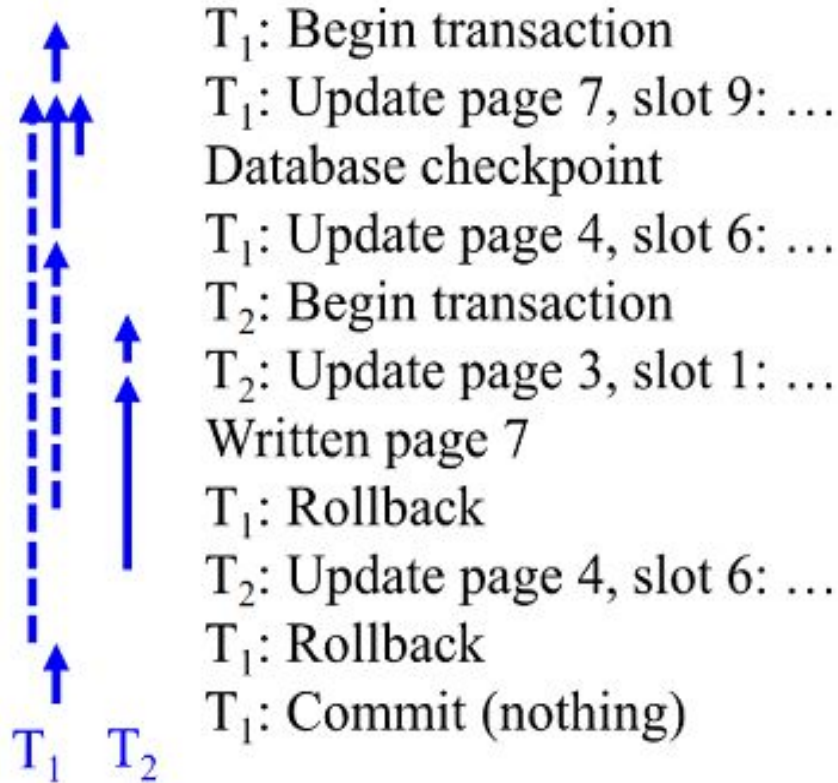
1990s: “compensation log records” in ARIES \Leftarrow row-level locking

- “Redo” only
- “Update back” vs restore page image
- “Commit nothing”, no end-of-transaction log flush
- Incremental lock release (during rollback)

2010s: “rollback log records”

- No data contents, no linked list per transaction

Per-transaction chains of log records?



Lock management during transaction rollback

1970s: simple & conservative

- Lock release at end-of-transaction

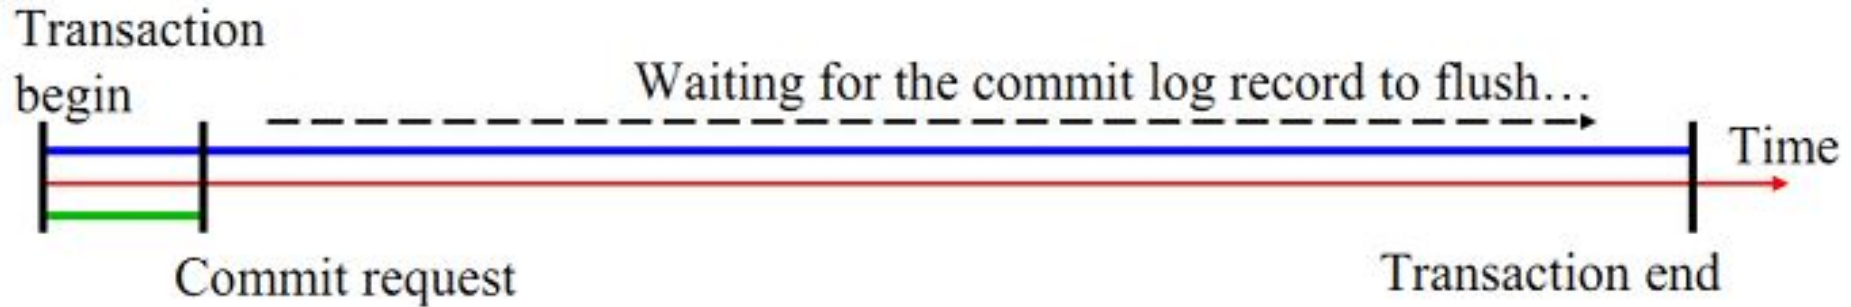
1990s: row-level locking [ARIES]

- Release of read-only locks \Leftarrow LSN for commit log record
- Incremental release of read-write locks (during rollback)
- “Save points”

2010s: controlled lock violation

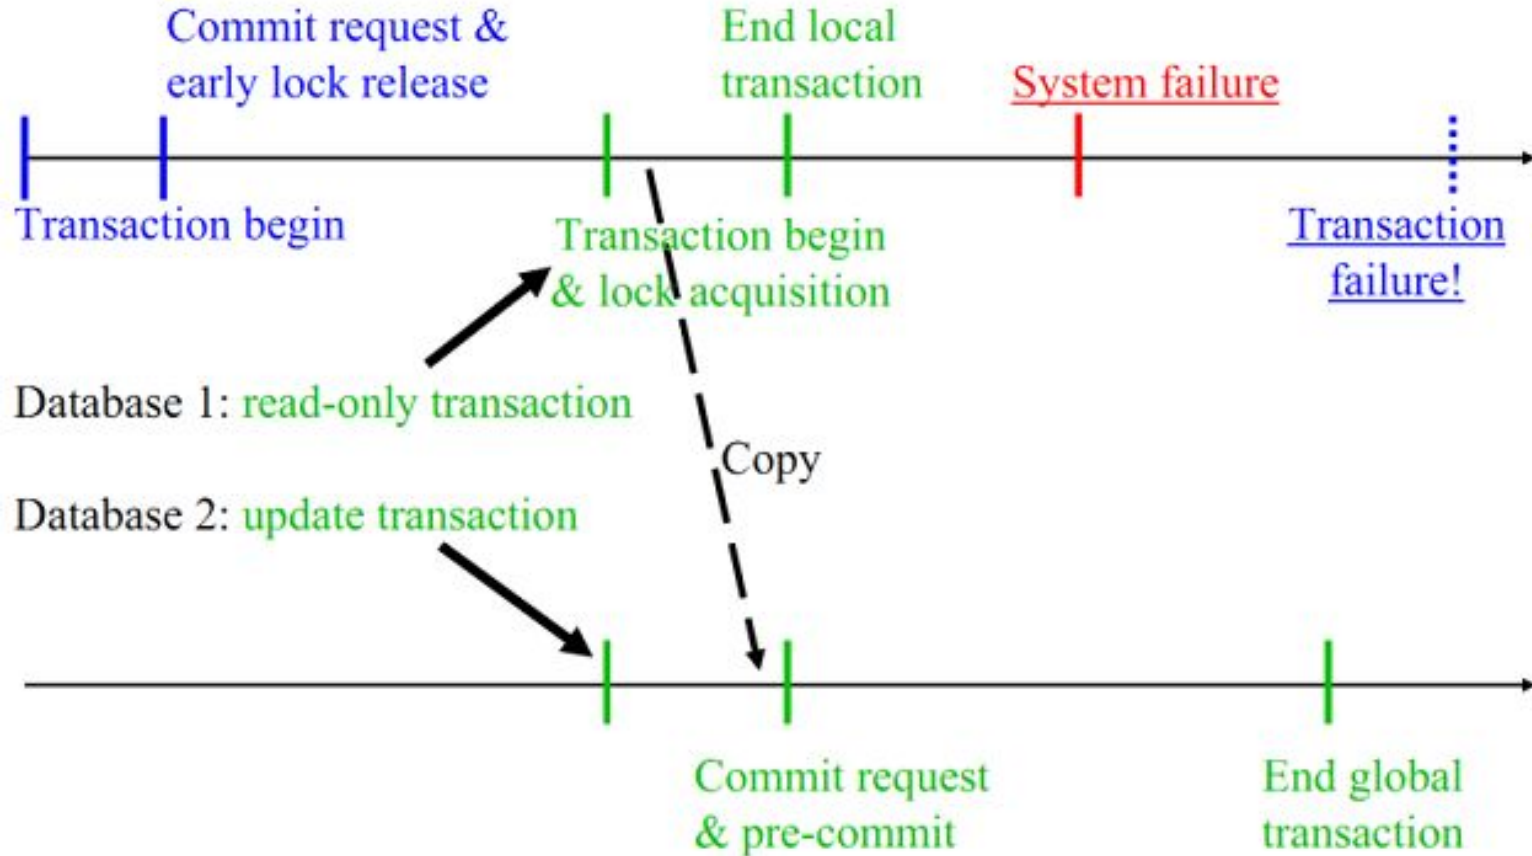
- Repair 1980s early lock release: controlled lock violation

Early lock release, controlled lock violation



Transaction phases → ↓ Techniques	Read phase = transaction & application logic	Commit logic			Hardening = force log to stable storage
		Commit preparation	Commit log record	Update propagation	
Traditional locking		n/a		n/a	
Controlled lock violation					

Premature publication with early lock release



Early lock release vs controlled lock violation

	Early lock release		Controlled lock violation
	S and X	S only	
Trigger	LSN for commit log record		
Performance	Best	Good	Almost best
Consequence of a write-read conflict	Premature publication	Waiting	Commit dependency

Commit delays due to controlled lock violation

Violator's scope → ↓ Violator's mode	Local transaction	Participant in a distributed transaction
Read-only in snapshot isolation	No locks, no lock violations, no delay	
Read-only with end- of-transaction commit	Delay up to flushing a commit log record	Delay up to flushing a pre-commit log record
Read-write	No delay beyond flushing a commit log record	No delay beyond flushing a pre-commit log record

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, **system**, media **failures**; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

System failure \Rightarrow system restart

Bad: server process & memory (transaction, lock, buffer pool managers)

Good: database, recovery log (even if slightly out-of-date)

1970s: System R



- Restart phases: log analysis, “redo” log scan, “undo” log scan
- Winner vs loser transactions

1990s: ARIES


- Lock re-acquisition during “redo”, new transactions during “undo”
- Transaction-by-transaction “undo”, no backward log scan

2010s: instant restart

Winner vs loser transactions



ARIES system restart and its log scans



Checkpoint information = current server state

For each active transaction:

- LSN of last update

Atomicity

For “undo” phase

with checkpoints

and new transactions:

Isolation

- Locks acquired and held
(ignoring read-only locks)

For each dirty buffer page:

- Page identifier
- LSN of the first update
(not yet on storage)

Durability

Availability

For instant recovery, replace:

- LSN of the last update
(instead of the first)

Classic restart: three phases & log scans

Log analysis \Rightarrow recover pre-crash server state

- Scan log from last good checkpoint to crash
- Identify transactions active at crash \Rightarrow “losers”

“Redo” \Rightarrow recover pre-crash database state (in buffer pool)

- Scan log from oldest unsaved change to crash
- Replay all missing changes in all database pages

“Undo” \Rightarrow transaction-consistent pre-crash snapshot

- Scan log from crash to oldest loser’s start-of-transaction
- Roll back loser transactions (write compensation log records)

Restart phases & their typical durations

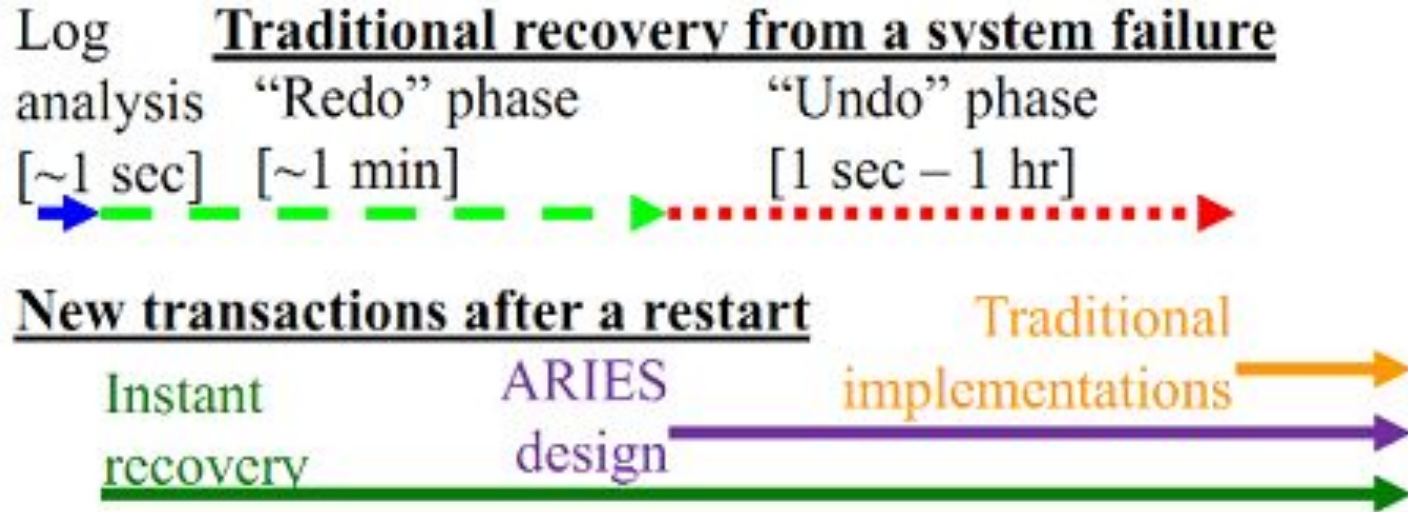
Classic design (1970s):



MS SQL Server improvements (as an example):

- 1998: new transactions after “undo”, parallel “undo” (tx rollback)
- ~2000: new transactions after “redo”, concurrent to “undo”
- ~2017: no “undo” phase, version clean-up like ghost clean-up
“constant-time recovery” = “accelerated database recovery”

Restart \Rightarrow new transactions, new checkpoints



\downarrow (apparent) mean time to repair $\Rightarrow \uparrow$ availability


Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, **media failures**; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Media failures \Rightarrow media restore

Bad: database (maybe only a single device)

Good: server process, log archive, recovery log



Optimizations:

- Remain online with database pages in the buffer pool
- All backups sorted on database page identifier \Rightarrow merge in single pass
- Sort the log records...

Point-in-time restore

From an earlier backup:

- Repeat missing history
- Roll back incomplete transactions

From a later backup:

- Roll back incomplete transactions
- Roll back transactions committed too late
- Roll back system transactions
- Alternative: roll back database history

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, media failures; **double failures**; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Multiple failures

Server crash during restart

- Re-start the same restart logic
- Early checkpoints!

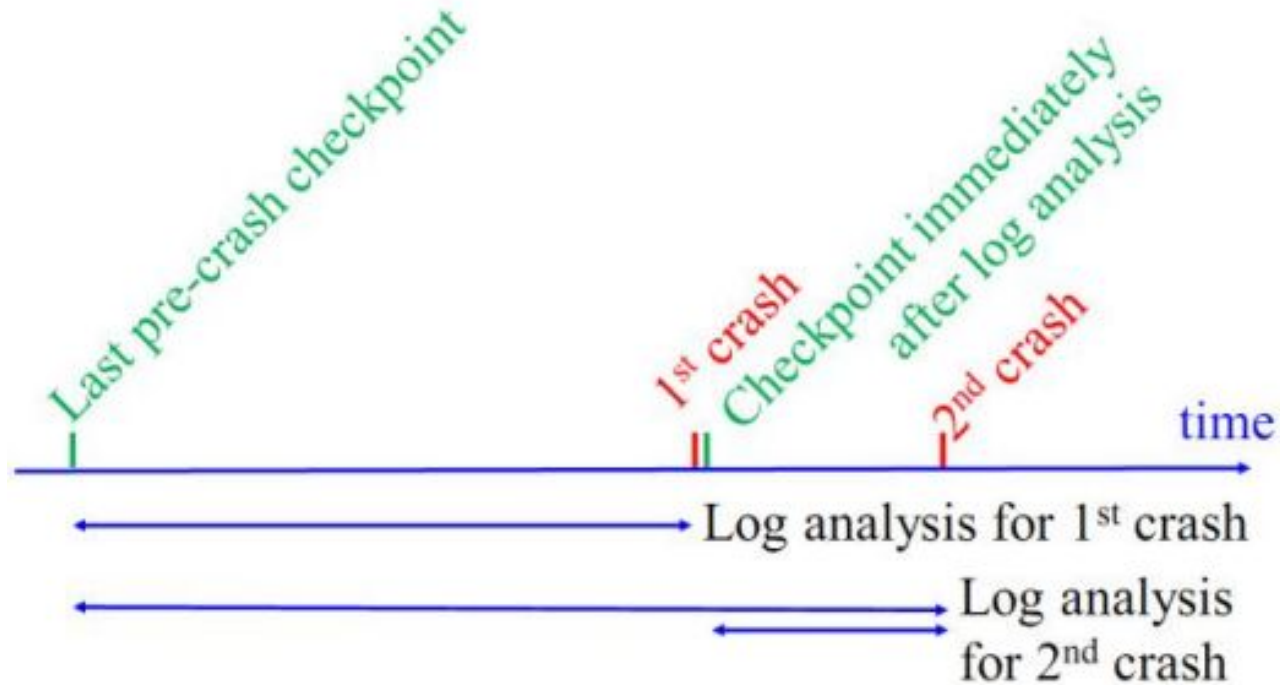
Media loss during restore

- Start over with new replacement media

System failure + media failure \Rightarrow reboot

- Restore, then restart
- Typical for in-memory databases and node failures


Double failure: checkpoint during restart



Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, media failures; double failures; **failover**
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Traditional log shipping



Node failure \Rightarrow node failover

Hot stand-by:

- Log shipping and continuous “redo” recovery
- Random I/O as much as the primary server

Cold stand-by:

- Capture backups + log archive
- Rely on restore + restart

Warm stand-by:

- Recent database image + up-to-date metadata + recovery log
- Incremental “redo”

Excursus: “log shipping to storage”


Amazon Aurora:

- Postgres or MySQL front-end
 - Local buffer pool, but no local storage
- “Storage” is another server
 - kept up-to-date by log shipping
 - serves buffer faults of the primary server
 - optional: shards & replicas

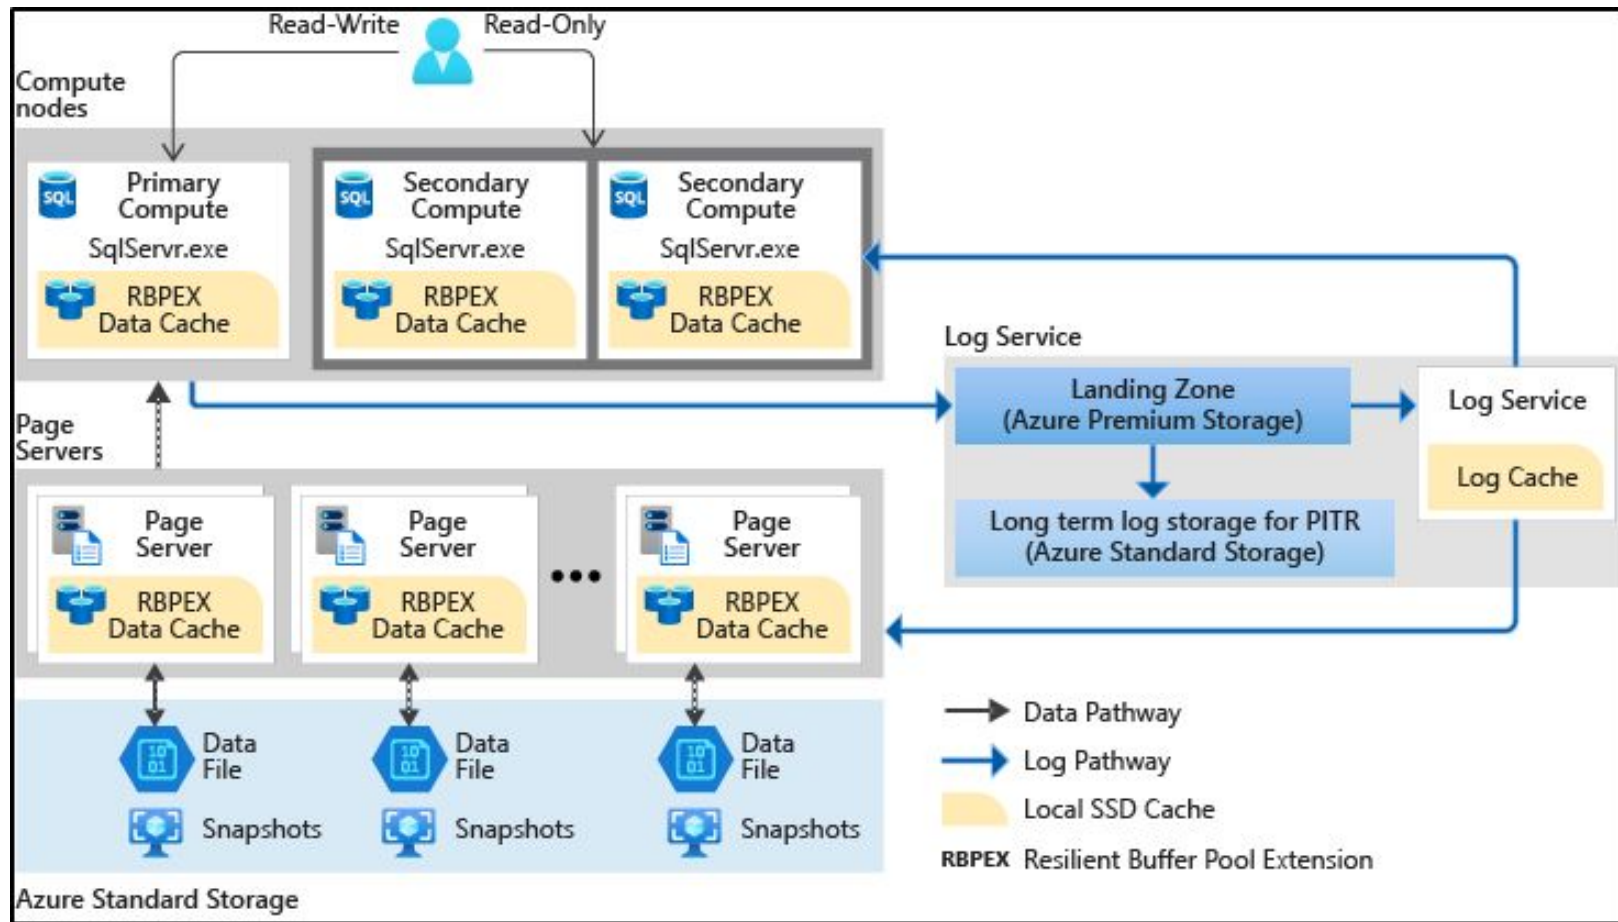
Assumption: storage imposes heavy CPU or memory load

- e.g., Microsoft Socrates = SQL Server hyperscale

“Log shipping to storage”



Microsoft SQL Hyperscale



Excursus: “log shipping to self”

- e.g., databases in persistent memory, Halloween protection
- Apply updates only after completed calculation of changes
 - Plan phase
 - Statement
 - Transaction
- ZZZ...

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- **Single-page failures**
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions


Failure scopes



Single-page failure



Media failure



System failure

SQL Server “online page-level restore”

1. Scan backups (full, differential, incremental)
extract most recent page image(s)
2. Scan log archive + recovery log
extract + apply “redo” for affected page(s)

Database locks & I/O: very little

Backup + log archive I/O: lots!

Single-page failure \Rightarrow single-page repair

Exploit per-page chains of log records

Anchor for the chain:

- Metadata
- Index parent node

Scale \rightarrow \downarrow History	Focused	Database-wide
Repeat	Single-page repair	“Redo” log scan
Rewind	Transaction rollback	“Undo” log scan

If fast enough, could be used deliberately – “write elision”, “read elision”

Per-page chains of log records

T₁: Begin transaction
T₁: Update page 7, slot 9: ...
Database checkpoint
T₁: Update page 4, slot 6: ...
T₂: Begin transaction
T₂: Update page 3, slot 1: ...
Written page 7
T₁: Rollback
T₂: Update page 4, slot 6: ...
T₁: Rollback
T₁: Commit (nothing)



MS SQL Server: consistency checks in “redo” pass, in log shipping, and in media recovery

Oracle: single-page rollback for snapshot isolation

Single-page failures: “Redo” recovery in reverse order since backup, move, or format

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: **self-repairing indexes**
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions


Single-page failure \Rightarrow single-page repair

Continuous comprehensive consistency checking

- Within nodes
- B-tree structure (pointers)
- Key ranges

B-tree structure



- Single pointer per node
- Fence keys



Foster b-trees, self-repairing indexes

Single pointer to each node at all times, local overflow gets “adopted”

ExpectedChildLSN with each parent-to-child pointer



Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - **Instant restart**, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Restart, new transactions, new checkpoints

Traditional recovery from a system failure



New transactions after a restart



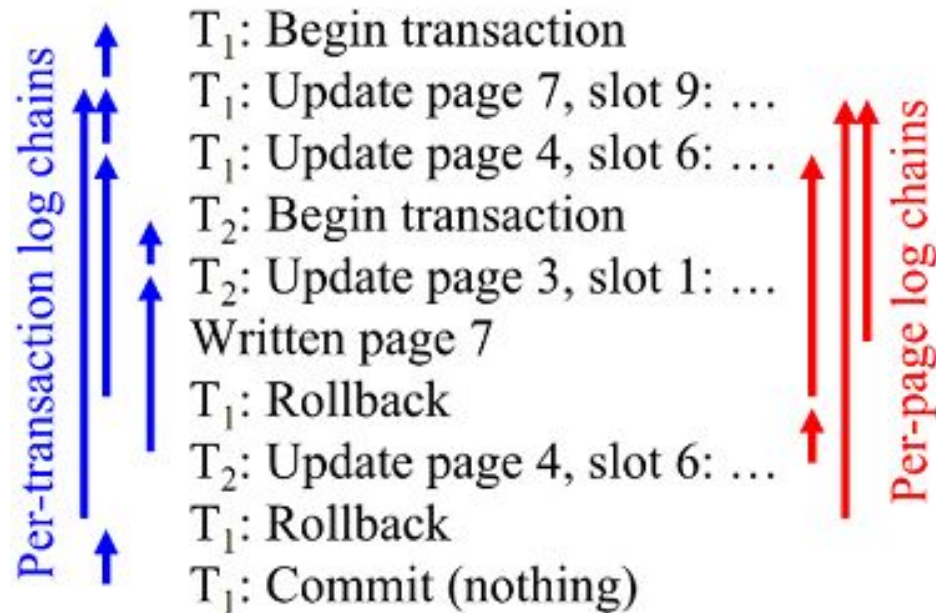
On-demand incremental “redo” and “undo”



Instant restart – before log analysis

Two transactions – interleaved

Server events – e.g., page written

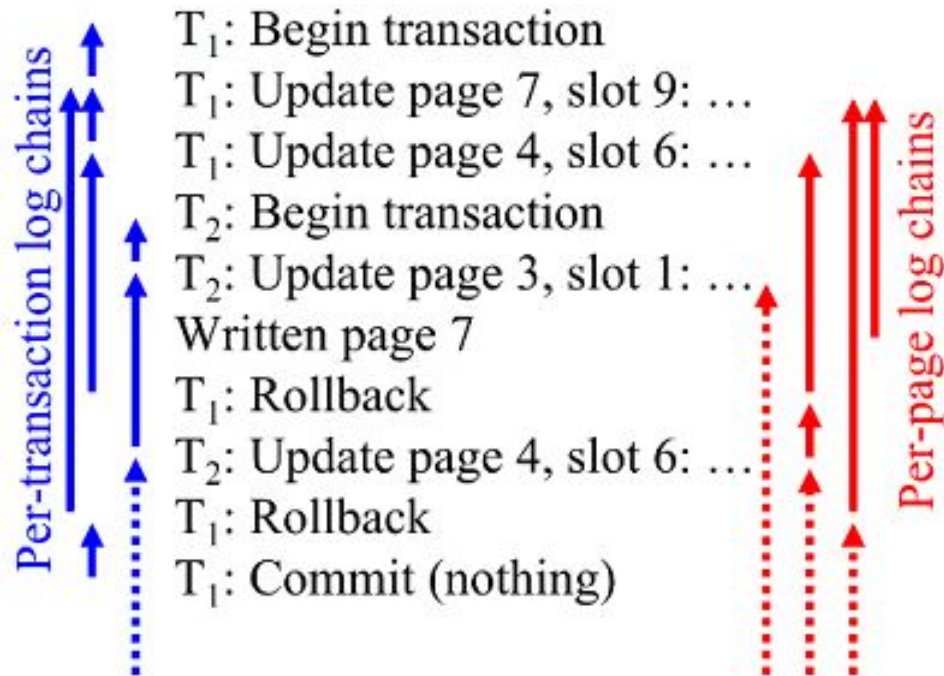


Instant restart – after log analysis

Log analysis recovers server state

- Transaction manager:
Loser transactions, LastLSN
- Lock manager:
Losers' read-write locks
- Buffer pool:
Dirty pages, LastLSN
(page descriptors only)

Log analysis enables checkpoints
and new user transactions!




Instant restart – on-demand “redo”

New transaction needs page 4

- triggers single-page repair
= incremental “redo”

Optimization opportunity:

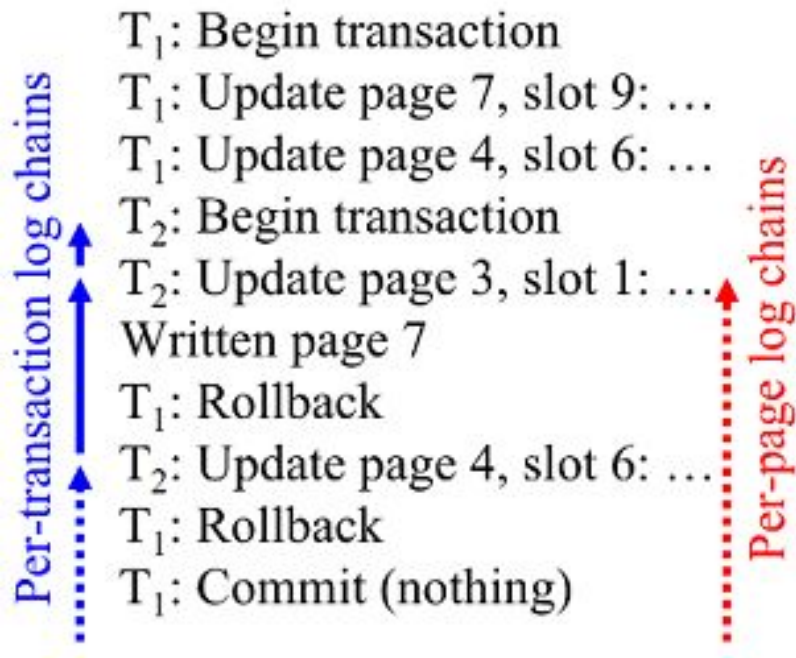
- Avoid two updates of slot 6:
rollback log record with
PriorLSNsamePage
== **PriorLSNsameTx**



Instant restart – on-demand “undo”

New transaction conflicts with
lock held by loser transaction T2

- triggers transaction rollback
= incremental “undo”
- can trigger single-page repair
= incremental “redo”



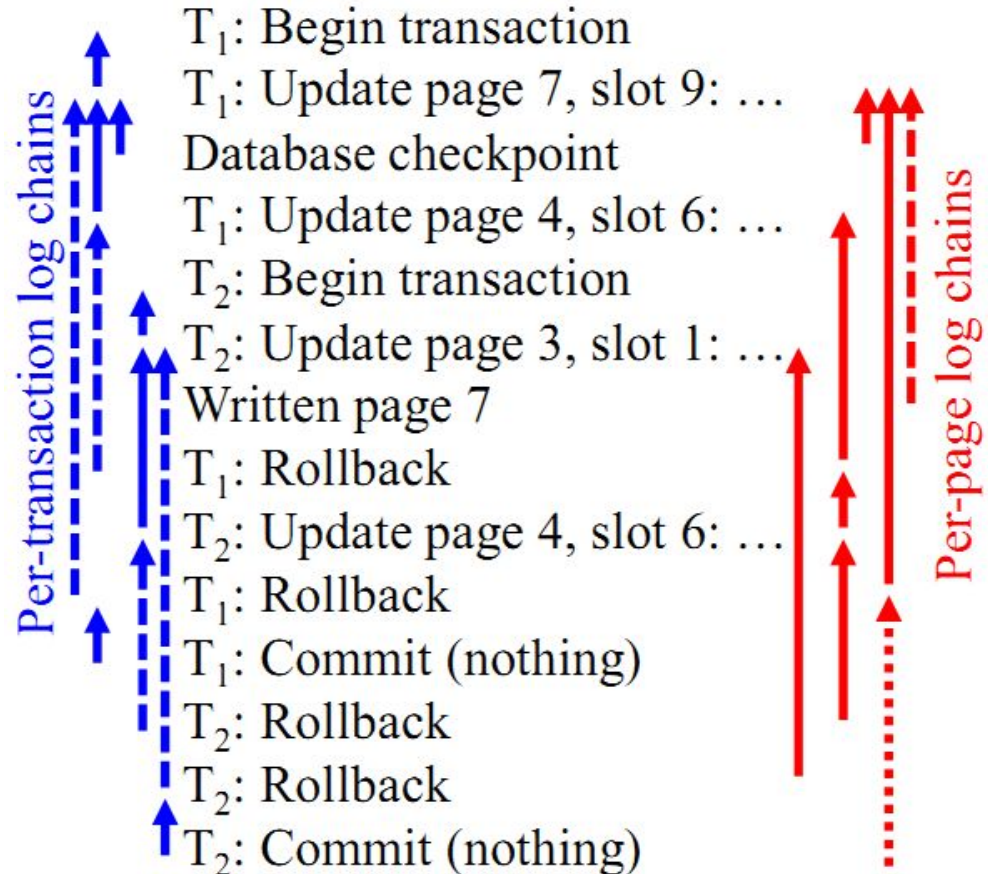
Instant restart – after transaction T2 rolls back

Pages 3 & 4 repaired

- page 7 still “in doubt”

Transaction T2 rolled back

- new rollback log records
- new commit log record



ARIES restart vs instant restart

	ARIES restart	Instant restart
Log analysis	Forward log scan	Forward or backward log scan Lock re-acquisition
“Redo”	Forward log scan + time-order log replay Lock re-acquisition	Single-page repair for each “in-doubt” page in the buffer pool – on demand and in many independent threads
“Undo”	Backward log scan + time-order rollback	Single-transaction rollback of each loser transaction – on demand and in many threads
New checkpoint	After “redo” phase	After log analysis
New transactions	After “undo” phase or after “redo” phase	
Concurrent transactions	A lock conflict blocks the new transaction	A lock conflict guides recovery during restart in order to unblock the new transaction


Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - **Single-phase restore**, instant backup
 - Instant failover, fail-back
- Conclusions

Single-phase restore

Sorted backups and log archive (partitions)


A single merge step!



Value-added in single-phase restore

A single pipeline!

Also: consistency checks of old and new backups + new database



ARIES restore vs single-phase restore

	ARIES restore	Single-pass restore
Transaction logging	Same techniques and costs	
Database backups	Offline or online full, differential, or incremental backups	Offline or online full backups
Log archiving	Copy and compress original recovery log	Partial sort + aggregation of log records
Restoring backups	One backup at a time; merge possible	Merge backup and runs of the log archive – random I/O with large transfers as in external merge sort
Log replay	In order of original execution – much random page I/O in the database	
Active transactions	Suspended during restore, rollback after restore is complete	
New transactions	Only after restore is complete	


Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, **instant backup**
 - Instant failover, fail-back
- Conclusions

Virtual backup


Backup = database + compression

Single-phase, just like single-phase restore



Remote virtual backup


Backup without network load, spike, provisioning



Instant backup

Backup + log archive = single-phase restore

⇒ equivalent to a fresh backup!



Agenda


- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - **Instant** restart, **restore**, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Instant restore

Single-page repair – incremental and on-demand

- Storage device divided into many (1K-1M) segments
- Database backup and log archive (partitions) sorted & indexed

Single-phase restore – efficient and in the background



ARIES restore vs instant restore

	ARIES restore	Instant restore
Transaction logging	Same techniques and costs	
Database backups	Offline or online full, differential, or incremental backups	Offline or online full backups, indexed by database page identifier
Log archiving	Copy and compress original recovery log	Partial sort + aggregation + indexing of log records
Restoring backups	One backup at a time; merge possible	Merge short runs from backup and runs of the log archive
Log replay	In order of original execution – much random page I/O in the database	
Active transactions	Suspended during restore, rollback after restore is complete	Online throughout, guiding restore operations
New transactions	Only after restore is complete	

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - **Instant** restart, restore, **reboot**
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Failure & recovery for in-memory databases

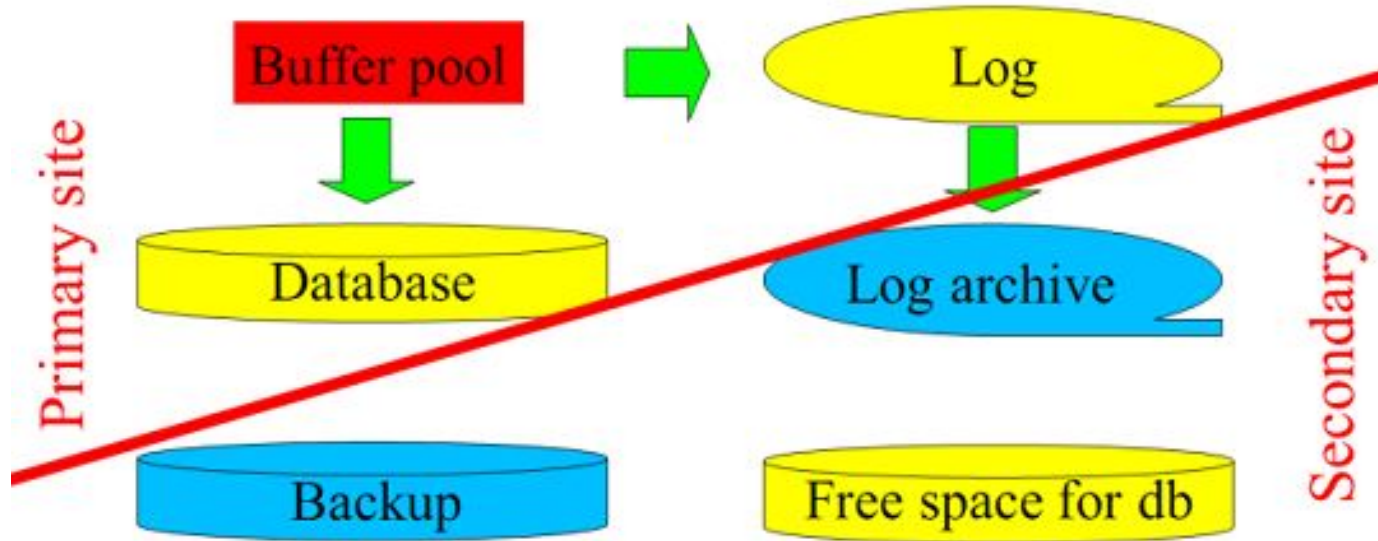
- System failure implies media failure
- **Reboot = restart + restore**
- Double failure \Rightarrow double recovery
 - Log analysis
 - Restore from backup + log archive: instant + single-phase
- New transactions & new checkpoints during recovery

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - **Instant failover**, fail-back
- Conclusions

Instant failover (to cold stand-by)

Instant restore + instant restart @ the secondary site



Comparison of techniques for high availability

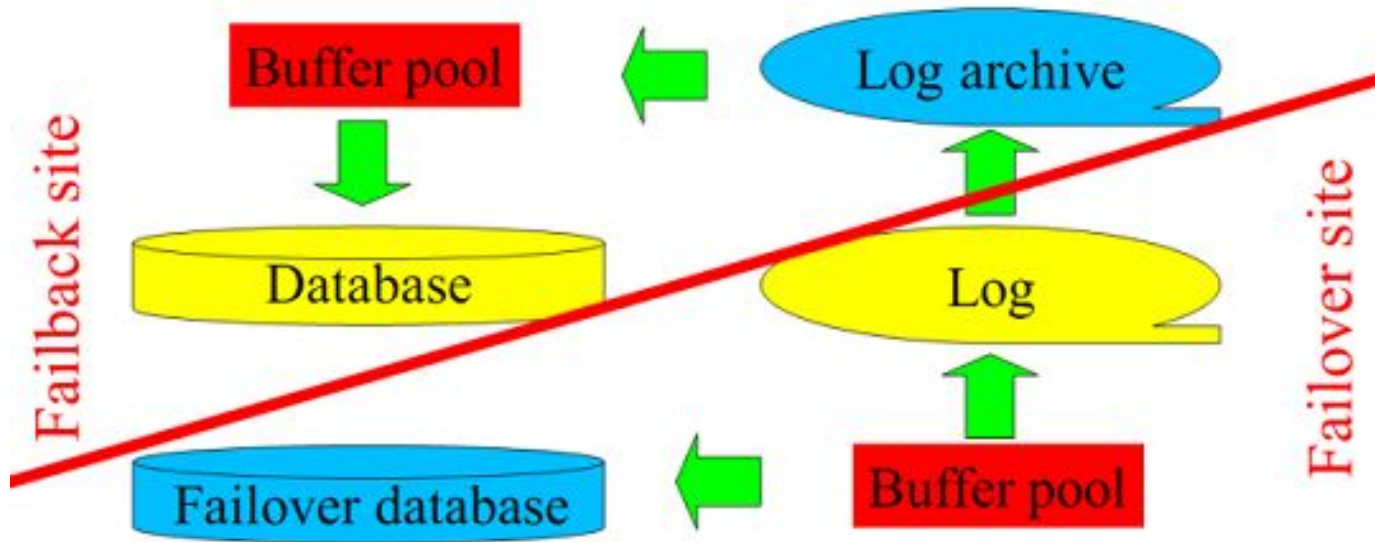
	Traditional mirroring	Traditional log shipping	Instant failover
Load on primary node before failure	Send entire pages	Send log records	
Load on secondary node before failure	Write received pages (random I/O)	Page updates: read, redo, write	Log archiving: sort + aggregate + index
Load after failover	Ship pages to create an additional mirror	Ship database backup and newer log records	Ship new log records
Vulnerability	Until an additional database copy exists and becomes up-to-date		Only during failover itself
Cost (plus backups, log, log archive)	Multiple up-to-date database copies		One up-to-date database copy
Per-node hardware requirements	Substantial CPU & I/O on each secondary node		Little CPU & I/O on secondary nodes
Network hardware requirements	Page copies, new backup	Log shipping, new backup	Log shipping

Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - **Instant** failover, **fail-back**
- Conclusions

Instant fail-back after repair of primary site

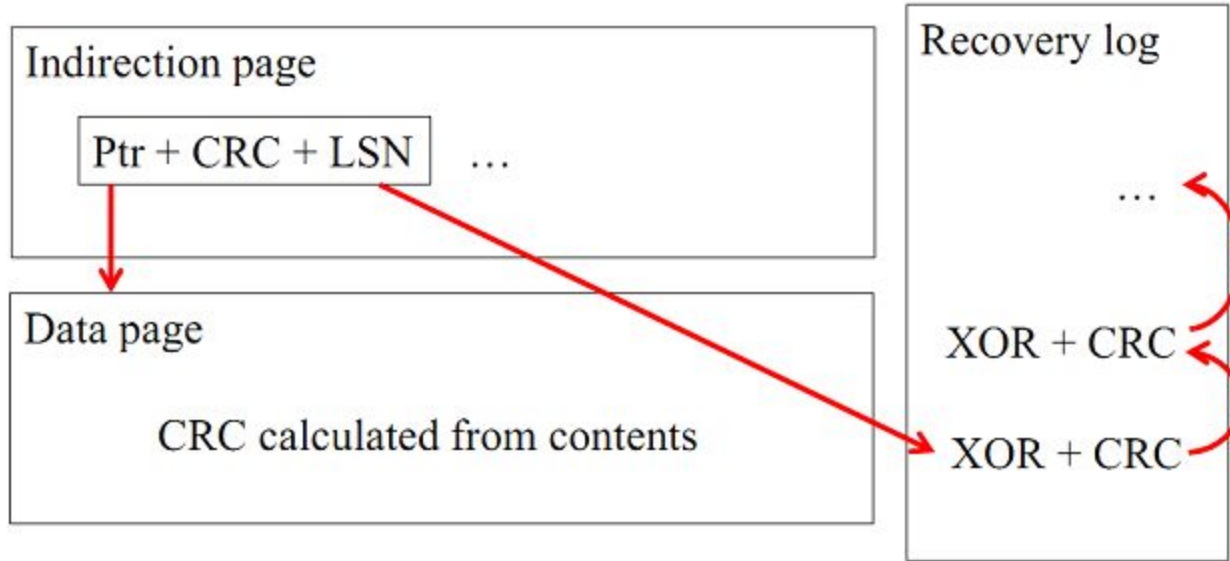
Instant restart @ the original primary site




Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- **Conclusions**


Write-ahead logging \Rightarrow self-repairing file contents




Direct archiving for large file contents




Performance of instant restart




Instant restart




ARIES restart – zoom



Instant restart – zoom



Performance of instant restore



Agenda

- Write-ahead logging
 - Commit, backups, archives, checkpoints
- Classic failures & recovery techniques
 - Transaction, system, and media failures; double failures; failover
- Single-page failures
 - On-demand local repair: self-repairing indexes
- Instant recovery
 - Instant restart, restore, reboot
 - Single-phase restore, instant backup
 - Instant failover, fail-back
- Conclusions

Thank you!