

DB101 – Storage formats

Goetz Graefe – Madison, Wis.

Agenda

- Introduction
- Devices, memory & storage hierarchy, caches & buffers
- Columnar storage, compression, zone filters
- Traditional row storage, primary & secondary indexes
- Domain indexes, master-detail clustering
- Log-structured merge-forests, stepped-merge forests
- Parallel & partitioned: mirrors, replicas, partitions, shards
- Free-space management, database catalogs
- Summary & conclusions

Topics omitted

1. Consistency checking – tables, indexes, databases, logs
2. Virtual, disaggregated, network-attached storage
3. Creation & maintenance algorithms, offline & online
4. Deferred maintenance of tables, views, and indexes
5. Automatic index tuning, ML @ DB
6. Physical database design as side effect of query execution
7. ...
8. Data quality and cleaning

Memory and storage hierarchy

- CPU caches: L1, L2, LL – “rules of 5×”
inclusive or exclusive, write-through or -back, LRU or ARC or ...
- Main memory: 500 cycles to serve a cache fault
traditional balance: 1 MIPS, 1 MB memory, 1 MB/s I/O; “5-min rule”
- Non-volatile memory: “imminent” for a decade
- SSD: 0.1 ms latency; flash translation layer
- Disk: 10 ms latency; 20 TB / 200 MB/s = 100,000 sec/scan
- RAID: 0: striped, 1: mirrored, 4: parity, 5: rotating, 6: double
- “Disaggregated” (network-attached & pooled) storage

Columnar storage

- Relation = rectangular table
- Queries retrieve & analyze specific columns in business intelligence, data analytics, “data wrangling” for ML

Speeding up large scans:

- Columnar storage
- Large transfers (I/O page)
- Parallelism
- Compression
- Zone filters

Compression in columnar storage

Compression feeds on repetition

- Storage efficiency (arrays without indirection)
- Dictionary encoding (best for strings)
- Frame-of-reference (best with correlation)
- Run-length encoding (best after sorting)
 - Equal values
 - Equal differences (e.g., in time series)

Zone filters

- Netezza: 3MB, single record type, separate min+max records
 - Infobright: bit vector filters
 - Zone filters: optional min+max, also second-to-min+max, bit vector
-
- bit vector filters: limited to exact-match look-up
 - min+max \Leftarrow correlation, e.g., order date & ship date
 - second-to-min+max \Leftarrow outliers

Row storage

- Page header + slot array (byte offsets within page)
PageLSN, record count
Poor man's normalized keys? Offset-value codes?
- Record header + slot array (byte offsets within record)
Field count(s), null bits?

In-page compression:

- In-page dictionary
- Variable-size integers (e.g., counts)
- In-string dynamic compression

“N-ary” in-page storage organization

0962	Jane	30
7658	John	45
3589	Jim	20
5523	Susan	52

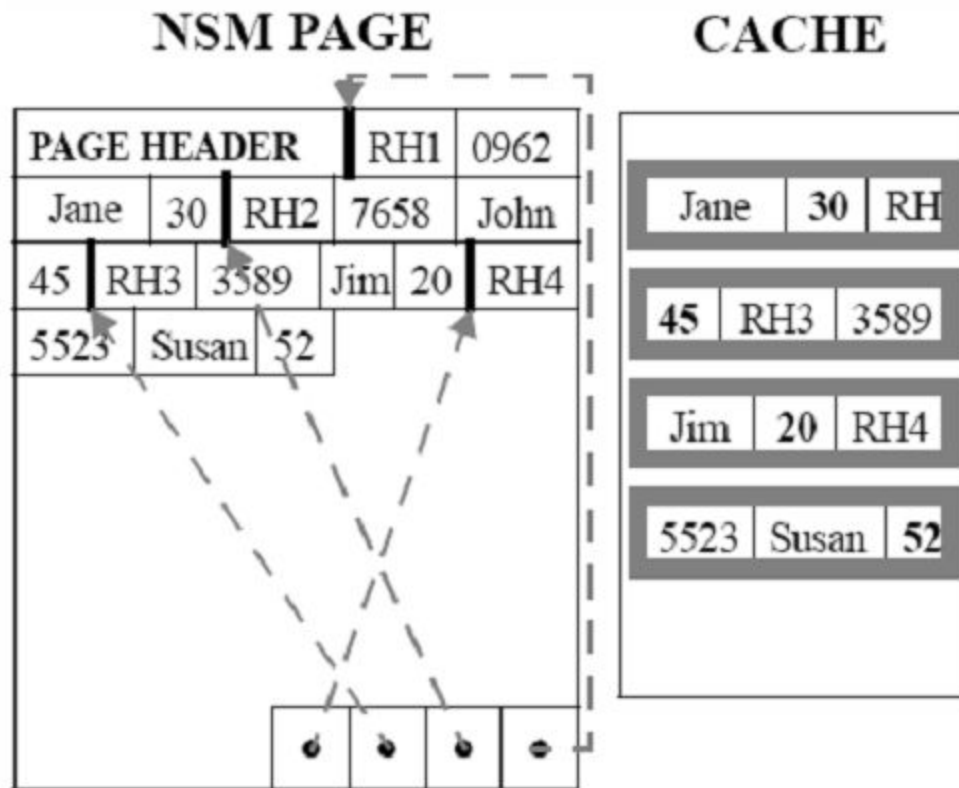


Figure 1. The cache behavior of NSM (from [ADH 01]).

PAX “partitioned across” in-page organization

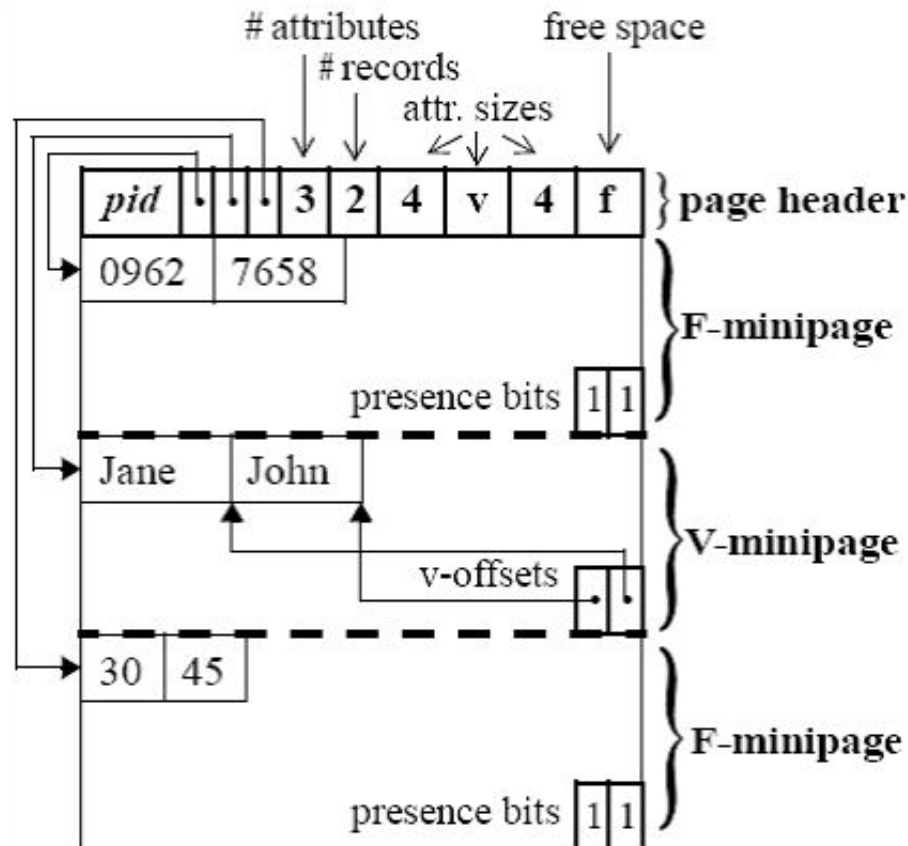


Figure 3. An example PAX page (from [ADH 01]).

PAX cache efficiency




Figure 2. The cache behavior of PAX (from [ADH 01]).

Indexes

- Map keys to values
C++ “map” templates, “key-value stores”
- Single-column index: value \rightarrow row identifier
- Compound index: (value, value...) \rightarrow row identifier
- Join index: value \rightarrow row identifier + row identifier
- Domain index: value \rightarrow table + row identifier
- ...
- Non-unique index: ... \rightarrow { ... }
- Compressed index:
- Bitmap index: { integer }

Databases vs key-value stores

Recall “sharing structured data:”

- No schema
 - No logical or physical data independence
 - No secondary indexes, no materialized views
- No privacy, security, or retention policies
- No query processing (only “get” and “put” in b-trees)
- No (general) transactions, e.g.:
 - No multi-row transactions, no cross-index consistency
 - No controlled redundancy, no distributed transactions
 - No phantom protection

Primary and secondary indexes

- Primary: all columns, defines row identifier
aka clustered index, index-organized table
- Secondary: key \rightarrow row identifier
 - Non-unique \Rightarrow multiple entries, { row identifier }, both
Bitmap = set of integers, hybrid, “word-aligned hybrid”
Compression of list vs bitmap
 - “Include” columns, e.g., all foreign keys
“Covering index” for “index-only retrieval”
Foreign keys: navigating index-to-index, record-to-record

Domain index

- e.g., $\text{CustId} \rightarrow \{ \text{tableId}, \{ \text{row identifier} \} \}$
all information about a customer with a single index search:
accounts, invoices, payments, shipments, returns, online reviews,
appointments, email, letters, phone calls, online activity,
loyalty program, referrals, ...
- Not offered much, not used much – not strictly relational
- Alternative: $(\text{CustId}, \text{TableId}) \rightarrow \{ \text{row identifier} \}$
- Refinement: merged index (multiple record types)

B-trees


- 1970, 1972: b-tree, b^* -tree, b^+ -tree...
- Today: data records in leaf nodes, “guide” keys in branch nodes
- Future: additional information with branch keys?
 - ExpectedPageLSN → self-repairing b-trees
 - “Small materialized aggregates” – efficient materialized views
 - Zone filters (on values, maximum difference, etc.)
 - ...

B-tree structure

- Creation: sort future index entries, build leaf-to-root leaves & nodes
- Index updates: sort future index entries, “merge into” b-tree
- O’Neil’s SB-tree (a b-tree within a b-tree)
- Write-optimized b-trees
 - Cheap page movement during write
 - Pointer swizzling in the buffer pool
 - Continuous comprehensive self testing – “fence” keys
- Foster b-trees – low latching requirements

O'Neil's SB-tree [1992]

- Allocate disk space in large blocks, allocate neighboring leaf nodes within blocks
- Split blocks 2→3 for ~85% disk utilization
- Fast index-order scans, e.g., large range queries





B-trees for correlated columns

- e.g., TPC-H lineitem ship date, receipt date, commit date
dates & timestamps often correlated in business intelligence
- In a b-tree index on “ship date” ...
... add min+max information on “receipt date” and “commit date”
⇒ reasonably efficient search on any date column
- ... add min+max of “transit days” and “spare days”
⇒ opportunities for compression with variable-size integers

B-trees on hash value and on Z-values

1. 0 dimensions \Rightarrow b-tree on hash values
index creation: sort on hash values
index intersection etc.: merge join on hash values
concurrency control: phantom protection on gaps between hash values
2. spatial, temporal, spatio-temporal \Rightarrow b-tree on space-filling curve
Z-values: bit interleaving
Peano, Hilbert, Moore, ... curves

Z-curve: bit interleaving



Log-structured merge-forests

- 1996/97: in-memory \Rightarrow fast, on-disk \Rightarrow slow
- Streams or bursts of insertions \Leftarrow
 - Update \rightarrow insertion of replacement record
 - Deletion \rightarrow insertion of “tombstone”/“anti-matter” record
 - Range deletion \rightarrow special anti-matter
- Alternative perspectives (all process insertions in memory):
 - “Repair” b-trees by spill+restart \Rightarrow binary merge steps
 - “External merge sort” b-tree entries \Rightarrow bursts of wide merge steps
 - “Never-ending sort” \Rightarrow space budget per merge level
 - Continuous merging \Rightarrow “merry-go-round” scans
 - Adaptive merging \Rightarrow rely on (impose) side effects of queries

LSM-forest as “write-optimized store”

- C-store/Vertica (HP, HPE, Microfocus)
read-optimized column store (no indexes)
write-optimized row store (in-memory index?)
- When sorting and merging for maximal compression, what is the extra cost (CPU, storage, data movement) for b-tree branch nodes?
- What is the achievable benefit?

Partitioning

- Orthogonal to “partitions” (runs) in log-structured merge-forest
- Focus on parallelism & manageability
- Partitioning per...
 - ... table (or multi-version) \Rightarrow “local” indexes \Rightarrow
local index-to-index navigation, local index intersection, etc.
 \Rightarrow little communication, high parallelism
 \Rightarrow good response time, bad resource consumption
 - ... index \Rightarrow “global” indexes \Rightarrow ...
- Co-partitioning on primary and foreign keys \Rightarrow efficient local joins

Partitioning

- Round-robin & random partitioning
- Range-partitioning \leftrightarrow distribution skew
- Hash-partitioning \leftrightarrow duplicate skew
- Hybrid partitioning: ranges of hash values
- Hybrid partitioning: hashing small key ranges
 - + distribution skew, + duplicate skew
 - Small query range \Rightarrow few parallel scans (1?), low latency
 - Large query range \Rightarrow many scans, high bandwidth

Free-space management

- Bitmaps with high concurrency
 - Per device or file (guides “backup” logic)
 - Per index (or partition) (guides “drop” logic)
- Manipulated in system transactions only
 - Model a b-tree with uncompressed bitmaps?
- Priority in restart and restore

Summary

- Devices, memory & storage hierarchy
- Columns for large scans (compression, parallelism, zone filters)
- Rows for index-to-index, record-to-record navigation